

Distributed File Systems


Case study : The Hadoop File System

Principles

- Open source implementation of the Google File System
- Distributed filesystem over multiple nodes
- Requirements
 - Fault tolerance
 - Scalability
 - Optimized for batch operations
 - Throughput more important than latency
 - Appending to files, no overwriting
- Provides a separate and global filesystem
 - Unix like paths
 - E.g. /home on HDFS is not /home of the machine
- Not part of the OS, added software
 - Written in Java, works on any machine with JVM support
 - Shipped with script for users and administrator

Architecture

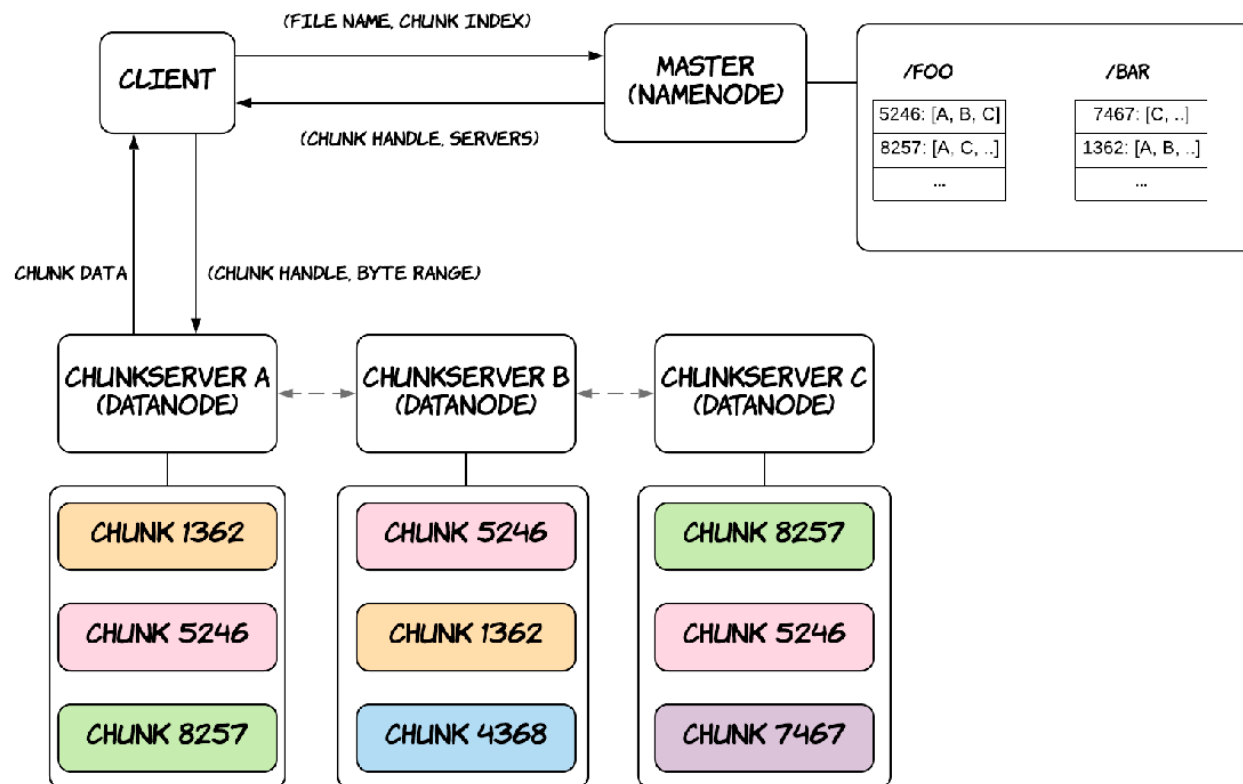
- Data and metadata are separated
 - Data are stored in **DataNodes**
 - Metadata are stored in **NameNodes**
- Data are divided into blocks
 - Default value 64MB or 128MB
- Blocks are stored on different machines



Architecture

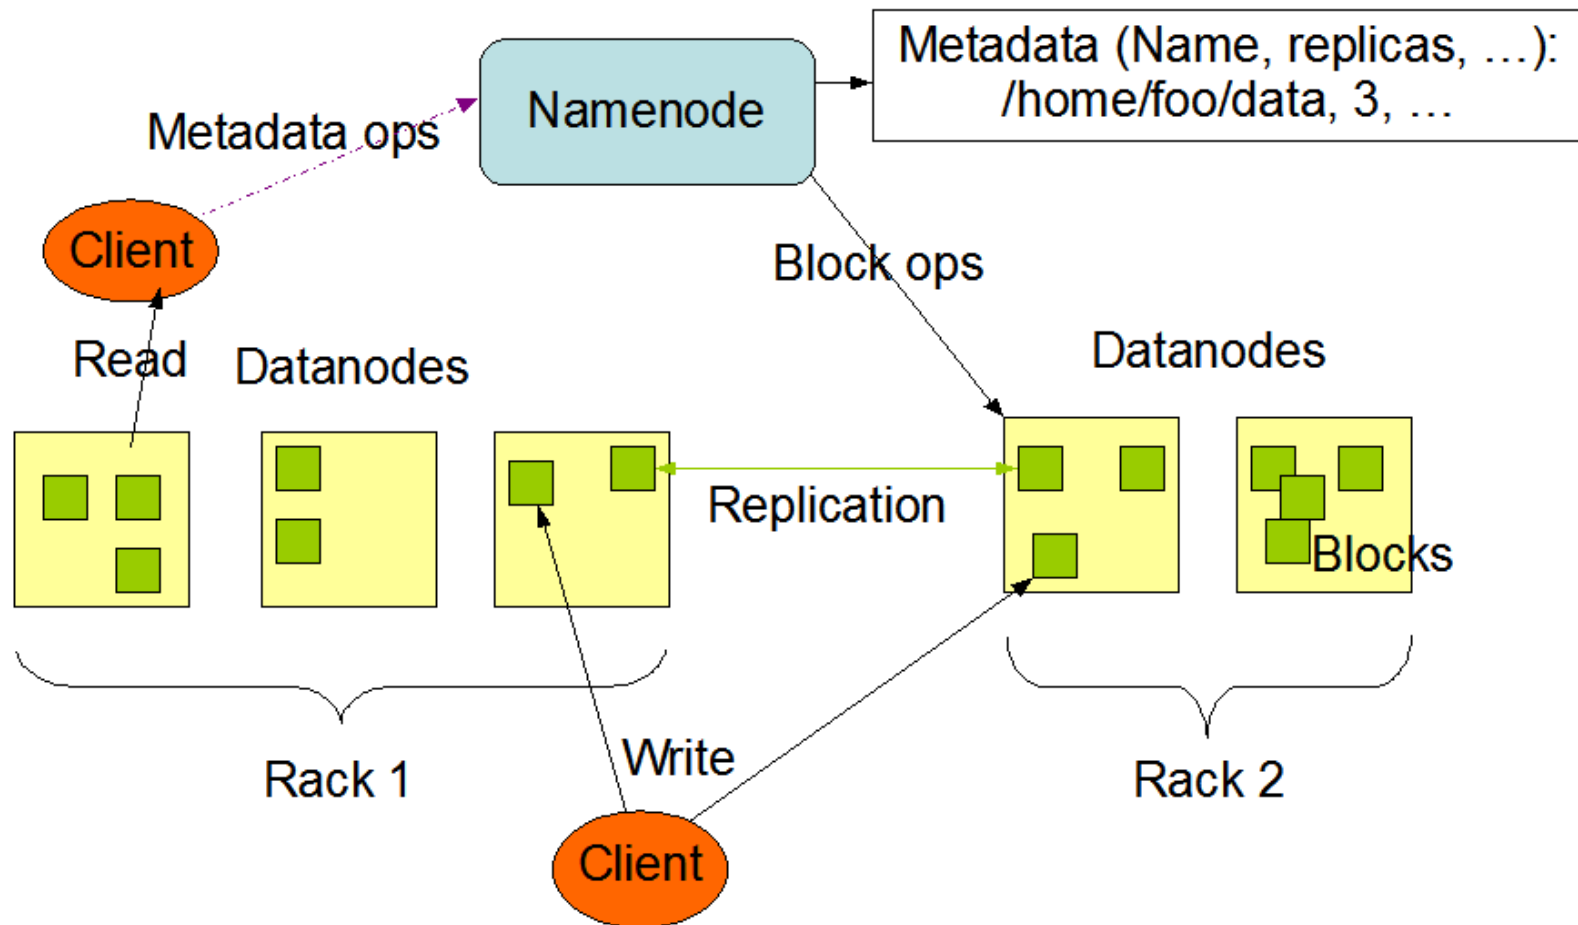
- Data can be replicated
 - Replication factor (default 3)
 - More costly (space) than RAID or erasure codes
- Not all operations are supported
 - No random write, only append and truncate
- Permission
 - User and group for coarse grained control
 - Set access to owner or group of users
 - Access Control Lists (ACLs) for finer control
 - Set specific permission to specific user
 - Disabled by default

GFS (HDFS) Architecture



Source: *Distributed Systems for Practitioners*, Dimos Raptis

HDFS Architecture



Reading & writing

- All clients read and write requests go through NameNodes
 - Validation and metadata
- Actual data go directly to clients
 - Faster access
 - Blocks can be requested in parallel
- Reading workflow
 - Client send request for file “/test.txt”
 - NameNode checks access rights and returns list of blocks+DataNodes

Reading & writing

- Writing workflow
 - Client contact NameNode
 - If writing allowed, check if file exists
 - If yes, error
 - NameNode returns a list of DataNode
 - Client sends data to DataNodes in round-robin
 - After writing all blocks, the client notifies the NameNode
- Replication is handled by DataNodes while receiving data

HDFS Blocks

- Large blocks
 - Not suited for storing small files
 - Limit overhead of metadata
- Transfer cost of a block
 - Disk access + latency + throughput
 - Large blocks minimize impact of disk access and latency
- Replication for
 - Fault tolerance
 - Tries to put replica on different machine and different racks
 - Faster access
 - Load blocks from node closer to client
- Also support Erasure Codes

Fault tolerance

- NameNode is a single point of failure
 - If down, cannot access data anymore
 - If destroyed (metadata lost), data are lost
- Possibility to use a Secondary NameNode
 - Maintains snapshot of metadata + edit log
 - Periodically apply edit log to metadata and store new state
- In case of crash
 - Restart NameNode
 - Get snapshot + log of Secondary NameNode and rebuilt recent state
- But
 - Rebuilding might be long
 - Might still lose some metadata

Fault tolerance

- *High Availability*
 - Feature introduced in Hadoop 2
- Support 2 NameNode
 - 1 active and 1 passive in standby
 - If active fails, standby takes its place
- How to ensure consistency?
 - Relies on JournalNode
 - Active write edit log to JournalNode
 - Passive regularly get edit log from JournalNode
- Protection against failures of JournalNode
 - Usually use 3 of them but N possible
 - Agreement based on a Quorum algorithm
 - Can tolerate $(N-1)/2$ failures

Commands

- All commands are done using bin/hdfs or bin/hadoop
 - `hdfs <command> [options]`
 - `hadoop <command> [options]`
- 3 types : client, admin and daemon
- Client commands : use the filesystem
 - *hdfs dfs <args>* or *hadoop fs <args>* : run the command args on the HDFS filesystem
 - Examples : -mkdir, -put, -copyfromlocal, -get, -copytolocal, -rmr

Commands

- Admin commands : manage the filesystem
 - `hdfs fsck` : check and repair the filesystem
- Daemon commands : manage the infrastructure
 - `hdfs balancer` : run the balancer utility
 - `hdfs datanode` : start a new datanode
 - `hdfs namenode` : start a new namenode