

# Processing Big Data

# Introduction

- How to efficiently process large amount of data?
  - Use many machines
  - Use many cores
- How to efficiently use many machine/cores
  - Use a suitable programming model
- What is a programming model
  - A way to write some program,
  - with access to a limited set of functions,
    - provided by libraries or frameworks
  - and an environment to execute it.

# MapReduce

# The MapReduce programming model

- Popularized by a paper from Google : **MapReduce: simplified data processing on large clusters (2008)**
- Simple model with 2 basic operations
  - Map
  - Reduce
- Assume data are structured as (key,value)
- Apply successive map and reduce operations
  - Not necessarily limited to 1 map and 1 reduce in theory

# Map and Reduce

- Map and reduce are functions with defined input-output

- Map :

- Uses a single (key, value) to produce multiple pairs
- Input : (key, value) or (key, \_)
- Output : One or many (key, value) pairs

- Reduce :


- Gets all values associated with a given key and produce new pairs
- Input : (key, [value1, value2, ..., valueN])
- Output : One or many (key, value) pairs

*can be different*


# Example

- Word count
  - Take a text, produce a list of (word, Nb occurrences)
- Map function :
  - Assume each word appears only once
  - Use word as key and add value 1
  - Input : (word, \_)
  - Output : (word, 1)
- Reduce function :
  - Sum all '1' for a given key (word)
  - Input : (word, [1,1,1,1,1,1])
  - Output : (word, 6)

## Map



## Reduce



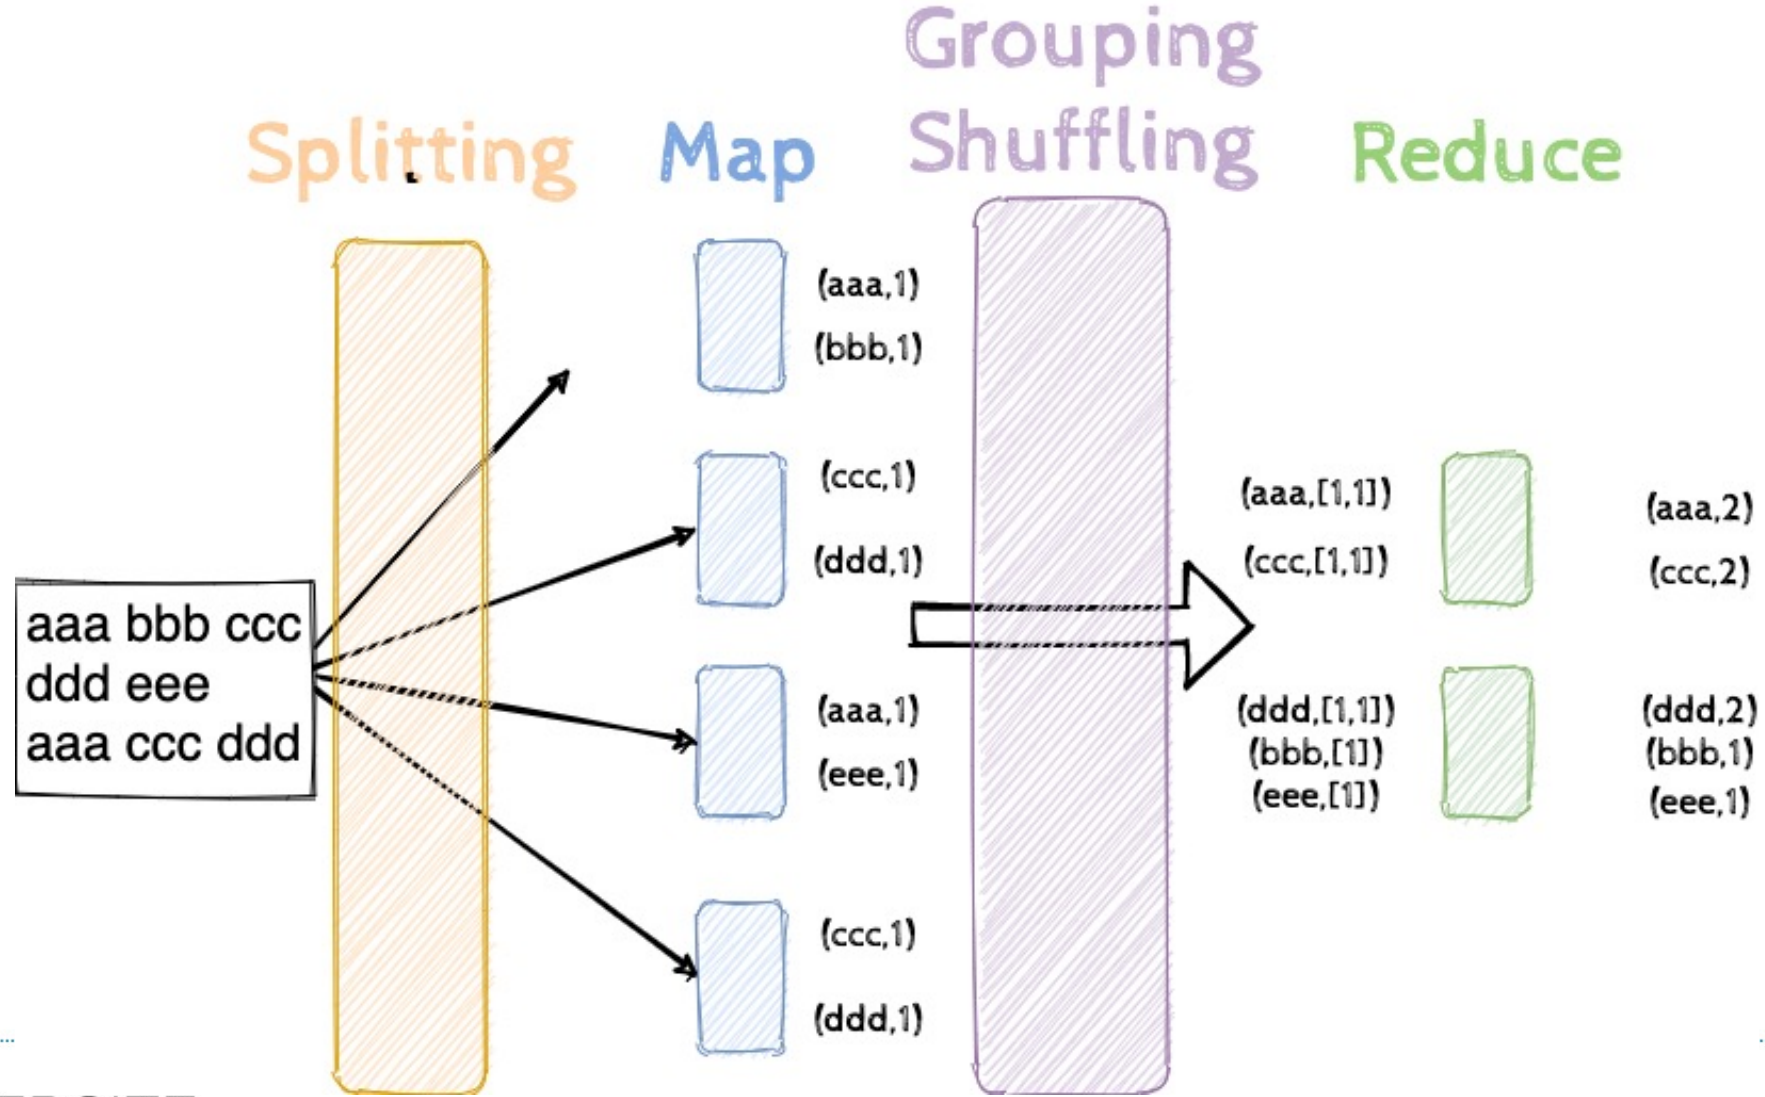
# Benefits of MapReduce

- A lot of real world problems can be expressed as MapReduce
- Maps and reduce functions can be executed in parallel
  - Map works on independent data
  - Reduce works on a single key
- Reduce can have inner parallelism
  - If complex, can reduce with multiple threads



# Implementations questions

- How is the input split into individual pairs for mappers?
- How are output of map grouped by key and sent to the correct reduce ?
- How is final result written ?
- All these are answered by a framework
  - All follow the same model but implementation may vary
- Basically, a MapReduce application has 4 phases
  - Splitting, Mapping, Grouping/Shuffling, Reducing

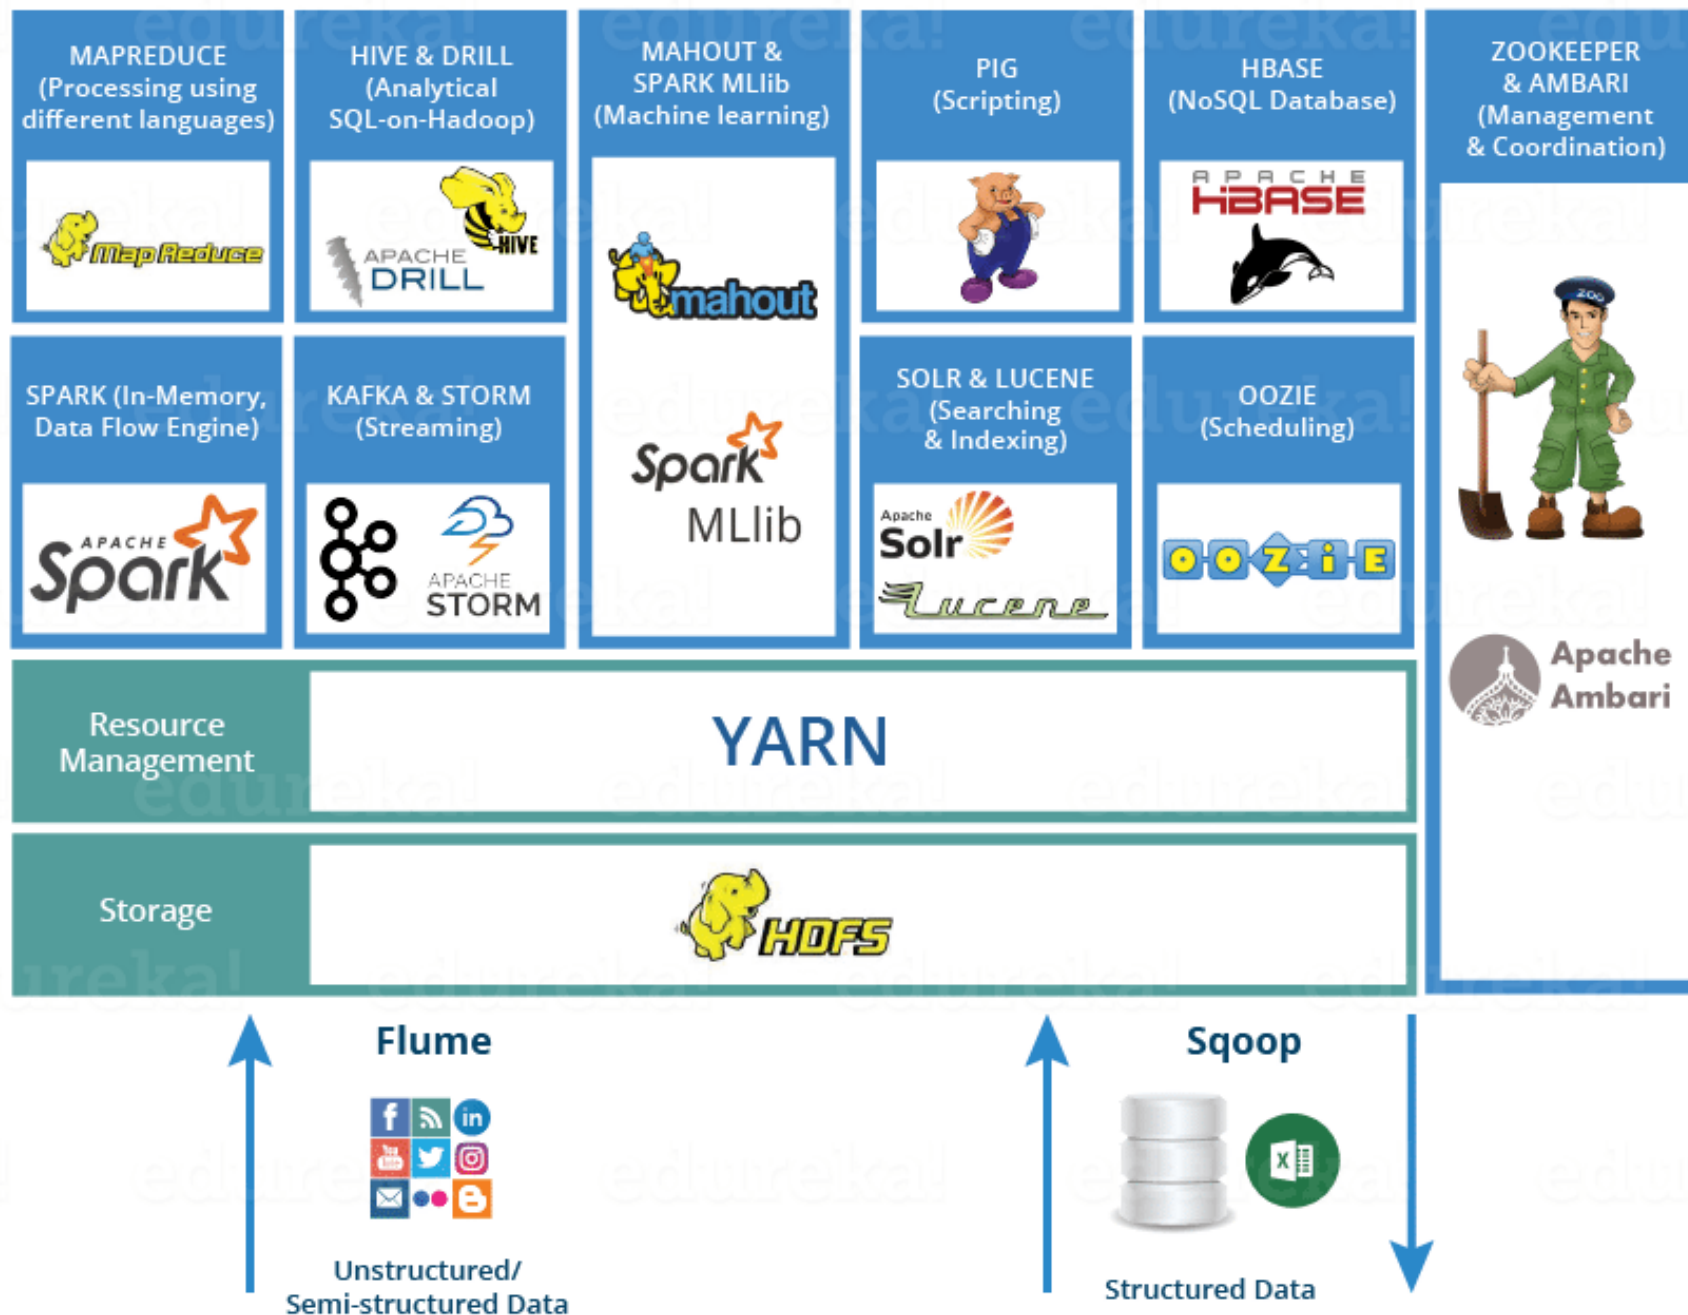


# MapReduce

The Hadoop Framework

# Introduction

- Hadoop is an open source implementation of Google MapReduce
- Provides full stack of services/frameworks
  - Not only MapReduce since version 2
- Most important components
  - HDFS
  - YARN
  - MapReduce



# YARN

- Yet Another Resource Negotiator
- Manages resources (nodes)
  - Knows the nodes available
  - Can provide nodes to an application
- YARN doesn't know or care about applications
  - So it can be used by any kind of application
- Main benefits
  - Nodes can be shared between user/applications
  - New applications/frameworks can use it and avoid managing resources
- Yarn is used outside of Hadoop

# Writing Hadoop MapReduce Programs

- A MapReduce program is called a Job
- Main language is Java
  - But can use any executable or script
- Map function implemented in a Mapper class
- Reduce function implemented in a Reducer class
- Default implementation
  - Splitting
    - If input is text file, key is offset, value is whole line
  - Shuffling
    - Based on hash value of key
- Package *org.hadoop.mapreduce*

(5280, "a b c")

# Datatypes

- Hadoop relies on its own serialization
  - More efficient than standard Java
  - Relies on the *Writable* interface
- Any key or value implements *Writable*
- Common ones
  - IntWritable : 32-bit Integer
  - LongWritable : 64-bit Integer
  - DoubleWritable : 64-bit Double
  - Text : String



# Writing a mapper

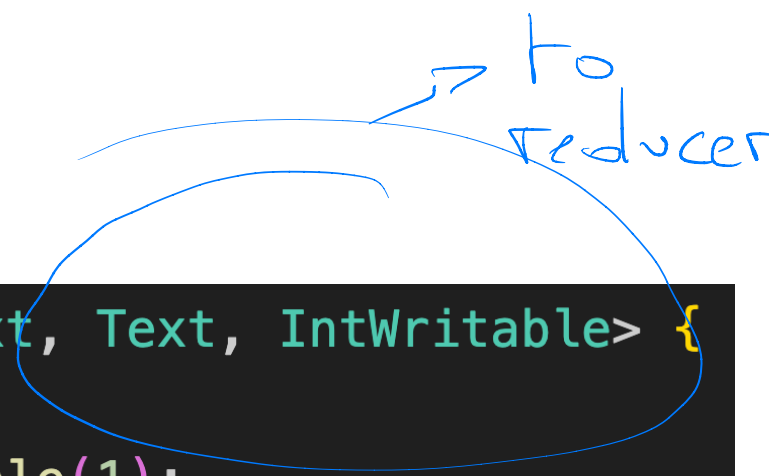
- A mapper takes *(key,value)* and produces *(key,value)*
- Steps for implementing your own mapper
  1. Decide on the data types of input and output
  2. Extends `Mapper<KEYIN, VALUEIN, KEYOUT, VALUEOUT>`
    - Set the types of all generics
  3. Implements *public void map(KEYIN, VALUEIN, Context)*
- Sending data to reducer is done through Context

# Writing a reducer

- A reducer takes *(key,[values])* and produces *(key,value)*
- Steps for implementing your own reducer
  1. Decide on the data types of input and output
    - Input comes from the mapper, output can be anything
  2. Extends `Reducer<KEYIN, VALUEIN, KEYOUT, VALUEOUT>`
    - Set the types of all generics
  3. Implements *public void reduce(KEYIN, Iterable<VALUEIN>, Context)*
- Output is done through Context

# Mapper example

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
    public void map(LongWritable key, Text value, Context context) throws... {  
        String line = value.toString();  
        //do something smart here  
        context.write(word, one);  
    }  
}
```



# Reducer Example

From Mapper

```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws... {  
        int sum = 0;  
        //do also something smart here  
        context.write(key, new IntWritable(sum));  
    }  
}
```

# Writing the rest and compiling

- In the *main(...)* method
  - Create a Job instance
  - Configure InputFiles, Mapper, Reducer, Output key and value...
- Everything has to be packaged as a jar
  - Use maven for dependencies and packaging

```
public static void main(String [] args) throws Exception
{
    Configuration c=new Configuration();
    String[] files=new GenericOptionsParser(c,args).getRemainingArgs();
    Path input=new Path(files[0]);
    Path output=new Path(files[1]);
    Job j=new Job(c,"wordcount");
    j.setJarByClass(WordCount.class);
    j.setMapperClass(MapForWordCount.class);
    j.setReducerClass(ReduceForWordCount.class);
    j.setOutputKeyClass(Text.class);
    j.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(j, input);
    FileOutputFormat.setOutputPath(j, output);
    System.exit(j.waitForCompletion(true)?0:1);
}
```

# Execution

- A Job is submitted to a Hadoop Cluster as jar
  - hadoop command
- All files have to be in HDFS
  - All paths relative to HDFS
- Number of mapper instances
  - Automatically decided based on the size (blocks) of input
- Number of reducers
  - Computed, can be set manually
  - Ideal value depends on the output of mappers

# Execution - 2

- Mappers/Reducers are executed close to data if possible
  - Use replication factor
- Result of a Job is written to HDFS
  - In a directory
- Output directory contains 1 file per reducer instance
  - Named *part-000xxx*
- Jobs cannot overwrite existing files
  - Remember to remove previous results before new execution

Error Launching job : Output directory `hdfs://localhost:9000/result` already exists



# Hadoop Streaming

- A simple tool to use almost anything as map and reduce functions
  - Part of the standard Hadoop distribution
- Mappers and Reducers can be any exec or script
  - Works with Python
- Mappers
  - Read from files on HDFS
  - Write to standard output as text with tab as (key value) separator
- Reducers
  - Read from standard input, assume tab as separator
  - Get (key value) pairs (!)
  - Write to standard output, automatically saved to file

# Hadoop Streaming

$(key, [val_1, val_2, val_3])$   
 $\{$   
     $(key, val_1)$   
     $(key, val_2)$   
     $(key, val_3)$   
 $\}$

- Limitations
  - No real grouping/shuffling
  - Reduce receives multiple (key,value) pairs
  - Relies on files and STDIN/STDOUT for data transfer
- Example :
  - `hadoop jar hadoop-streaming-3.1.4.jar -input /sample-text-file.txt -output /results -mapper mapper.py -reducer reducer.py`