

Processing Fast Data

Context and motivation

- So far, we have seen how to store and analyze data
 - Assume all data is available
 - Process whole dataset or restrict it manually (i.e. in algorithm)
- Fine if
 - It makes sense
 - Latency is not the main factor (e.g. a couple of minutes to get results is fine)
- But what if
 - We want the result fast
 - There is no notion of “whole” dataset, new data is created all the time

Data streams

- A stream is an unbounded sequence of elements
 - No end to a stream
 - New data becomes available over time
 - Some might become unavailable or outdated
- Elements of a stream have timestamps
 - Generation time as measured by the source
 - Arrival in the system as measured by the system
- Elements are produced by external sources
 - No control on the volume and the arrival rate
- Systems specializes in stream-processing are called Data Stream Processing (DSP) systems or Stream Processing Engines (SPE)

The 8 requirements of Real-Time Stream Processing

- *Michael Stonebraker, Uğur Çetintemel, and Stan Zdonik. 2005. The 8 requirements of real-time stream processing. SIGMOD Rec. 34, 4 (December 2005), 42–47.*
- To efficiently process data-streams, a system must have these features
 - 1) Keep the data moving
 - Data should not be stored for processing
 - 2) Support a high-level, stream-oriented language
 - Need for adapted primitives and operators

The 8 requirements of Real-Time Stream Processing

3) Provide resilience against streams imperfections

- Missing or out-of-order data

4) Be deterministic

- Outcome should be predictable and repeatable

5) Have an efficient way to store, access or modify state information

- This information can represent results on past events
- Allows processing new data with past information
- Should not add latency (no complex or remote database access)

The 8 requirements of Real-Time Stream Processing

6) Ensure availability and data integrity

- Use high-availability solutions like master with failover
- No data should be lost

7) The system should be scalable

- Distributed with automatic horizontal scaling

8) Process data immediately

- Low overhead and low latency

General principles

Models

Processing models


- Batch processing
 - Wait for all data before processing
 - High latency but high efficiency
- Event-based processing
 - Process each new data as they arrive
 - Low latency but high overhead
- Micro-batching
 - Process new data in small batches
 - Try to balance latency and overhead

Programming model

- Dataflow (or datastream) programming
 - Program is a Directed (Acyclic) Graph or D(A)G
 - Usually called topology
- Each vertex is an operator
 - Data flow through edges
- Data usually have
 - Timestamp
 - Key-value format
 - Sometimes called tuples

Example : counting specific words


topology



Programming model

- Assume a shared-nothing architecture
 - Operators can have internal state
 - Work on separate streams/partitions
 - Avoid global state in the system
- Easy parallelism and distribution
 - Split data among operators
 - Scaling-out operators

Example : counting specific words



Execution Model

- Exactly ~~one~~ *ONCE*
 - Each data is processed exactly once
 - Intuitive model, but hard to enforce
- At least once
 - All data will be processed once
 - Some might be processed multiple times

General principles

Operators and topologies

Common operators


- Some operators appear often in apps/papers
- Stateless operators
 - Map, filter,
- Stateful operators
 - Work on multiple data, including previous
 - Sum, count, min, max, average

upstream
←


downstream
→

Classical topologies


Linear



Diamond



Star



UNIVERSITÉ
CÔTE D'IVOIRE


• Upstream vs Downstream

Windowing

- Sometimes need to process elements in group
 - But a datastream is unbounded!
 - Need to artificially create groups
- Window
 - A finite set of elements to be processed together
 - Can cover multiple micro-batches
- Elements can be grouped by
 - Timestamp (emission, arrival, processing...)
 - Count
 - Other criteria


Windowing

- Tumbling
 - No common elements between successive windows



Windowing

- Sliding
 - Windows can overlap




General principles

Performance

Metrics

- Processing time
 - Measured at each operator
- Latency
 - Duration between arrival of data and its processing
 - Can be measured at various points
 - End-to-end latency indicates how reactive the system is
- Throughput
 - Number of tuples processed per time-unit
- Latency and throughput can change independently
 - Example ?



Backpressure

- Not all operator of a topology have the same speed
 - Consider slowest operator in DAG
- How to manage slow operators?
 - Put incoming tuples into a queue for buffering
 - Risk of ever-increasing queue
- Solution
 - Notify upstream to slow down
- Backpressure helps with temporary overload

Case Study : Apache Storm


Introduction and concepts

Introduction

- A distributed Stream Processing Engines
- Written in Clojure
 - Dialect of Lisp on the JVM
- Bought by Twitter, open sourced
- Now an Apache project

Heo U

Basic concepts



- 2 basic components
 - Spouts
 - Bolts
- Data are Tuples
 - Travel on Streams
- Interconnected to form a topology

Basic concepts

- Tuples
 - Basic data element
 - Contains named fields
 - Contain integers, longs, shorts, bytes, strings, doubles, floats, booleans, byte arrays or any Serializable type
- Streams
 - Unbounded sequence of tuples
 - Streams have an id
- Spouts
 - Data source
 - Read data from external sources (file, network,...) and emit them on Streams

Basic concepts

- Bolts
 - Data processing unit
 - Receives data from Streams, treat them and send them to possibly other bolts
- Tasks
 - Replicas of Spouts and Bolts
- Topology
 - An interconnection of spouts and bolts
 - Direct Acyclic Graph
- A Storm application
 - Contains a topology
 - Runs indefinitely

no replicas
id=1 → 0

Stream grouping

- Each bolt must specify which stream(s) it will receive
 - How is the stream partitioned if multiple replicas ?
 - Need to choose a strategy, called Stream Grouping
 - 8 of them in Storm
- Shuffle Grouping :
 - Default, random distribution, with guaranteed equal load
- Fields Grouping
 - Stream is partitioned by fields from Tuples
 - Tuples with the same value for a field are sent to the same Task
- All Grouping
 - The stream is replicated to all Tasks
- Global Grouping
 - The stream goes to the Task with the lowest id

3 replicas

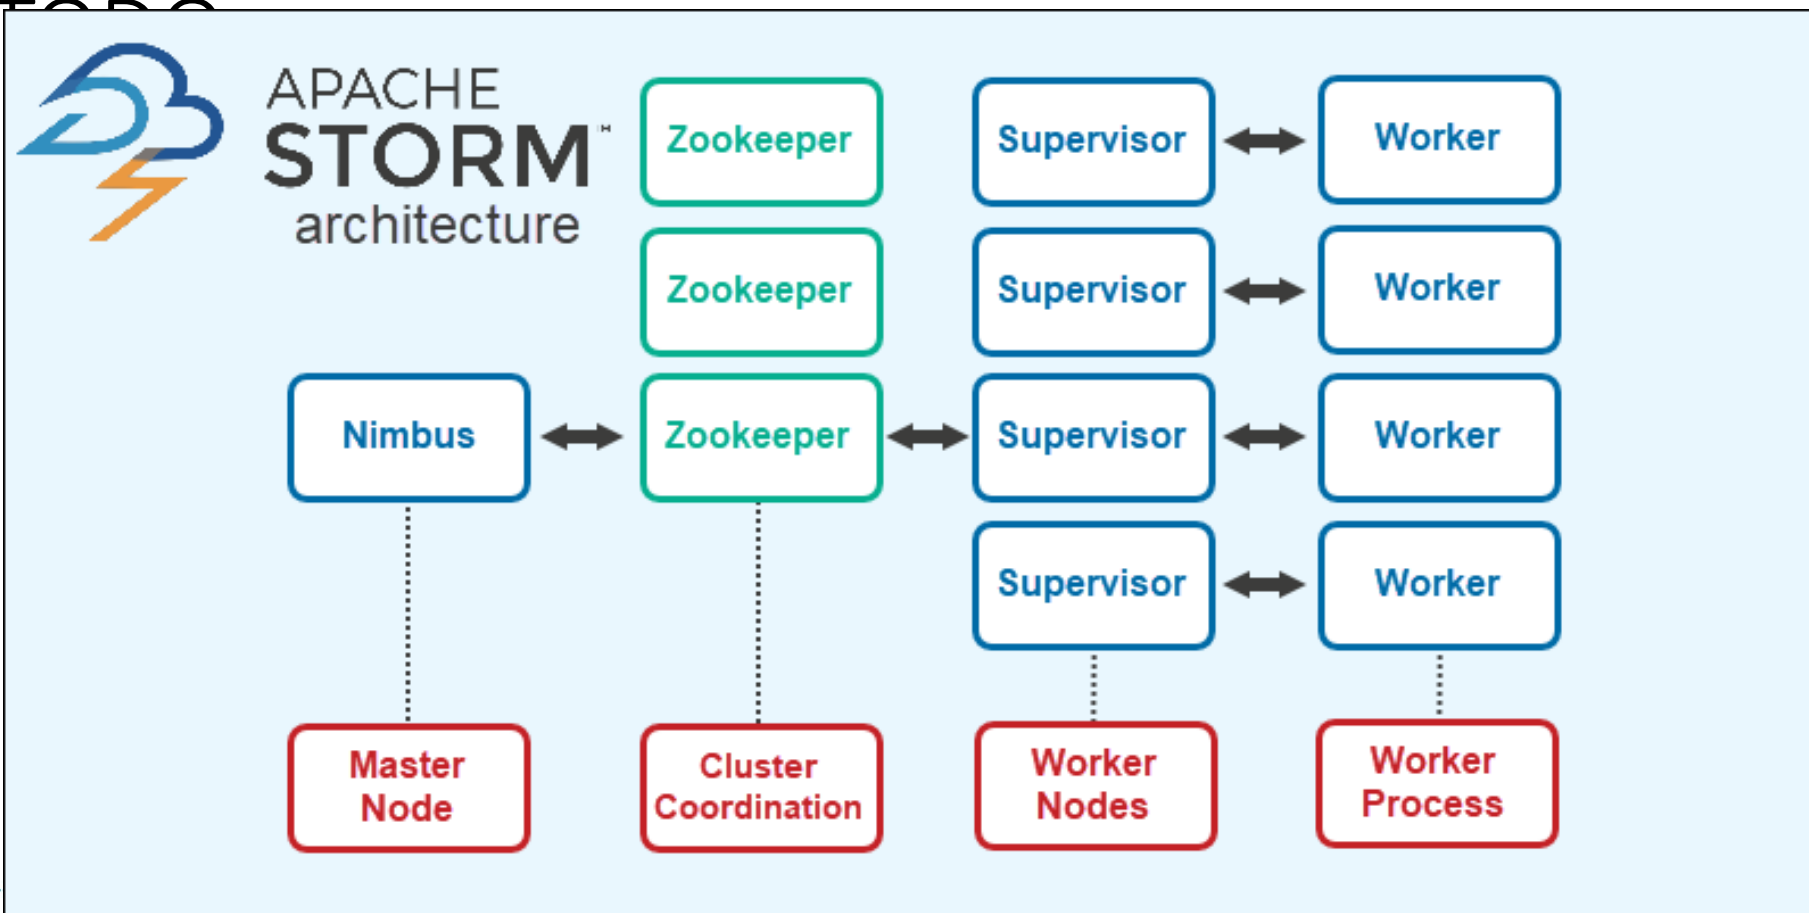
id=1 → 0
0
0
0

Architecture

- Distributed architecture
- *Nimbus*
 - Controller Node
 - Manage new topologies, fault-tolerance, scheduling
- *Zookeeper*
 - Manage coordination between all components and cluster resources
 - Not a Storm specific component
- *Supervisor*
 - Provide *worker* (JVM) to *Nimbus*
- *Workers*

→ Seure

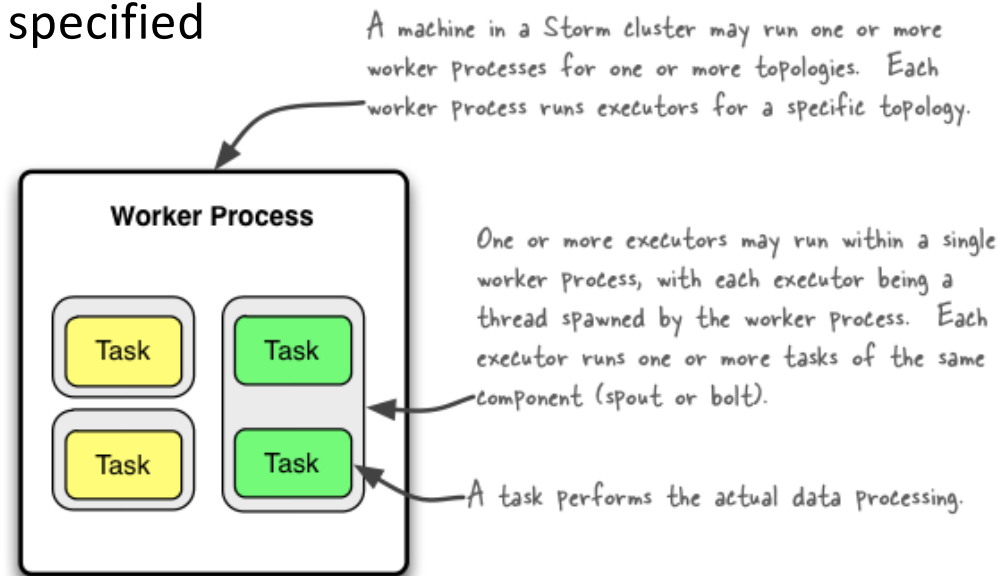
TOPIC



Worker

- *Worker*

- A JVM containing *Executors* (threads)
- Each worker is assigned a communication port (TCP), used also as ID
- Executors process *tasks* (*Spouts and Bolts*)
- Number of executors can be specified
 - *Parallelism hint*

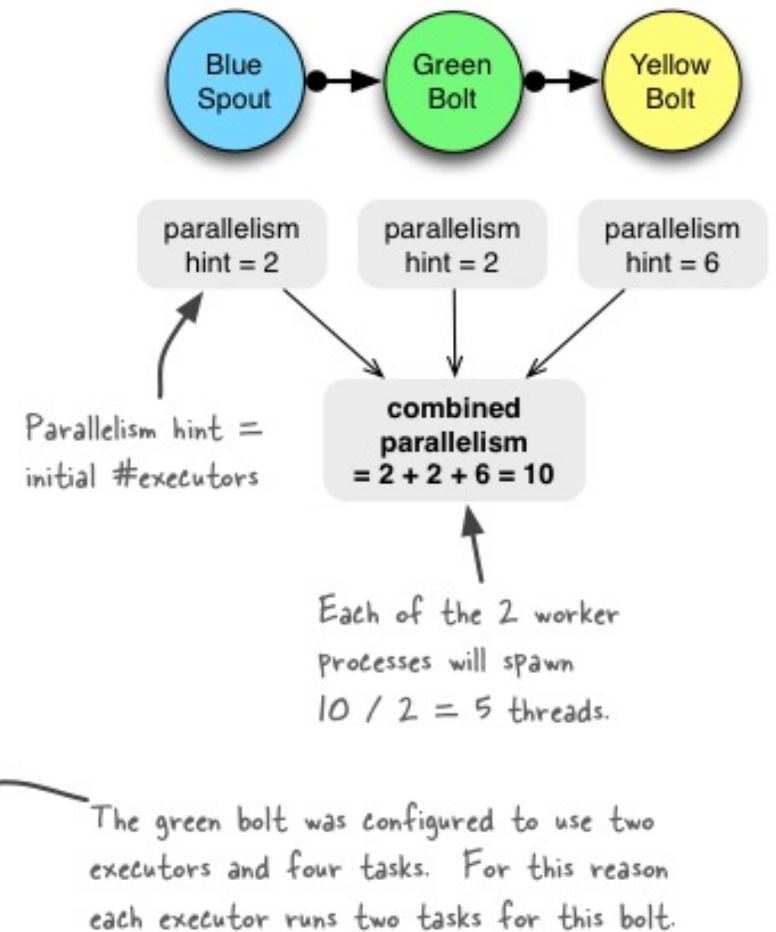
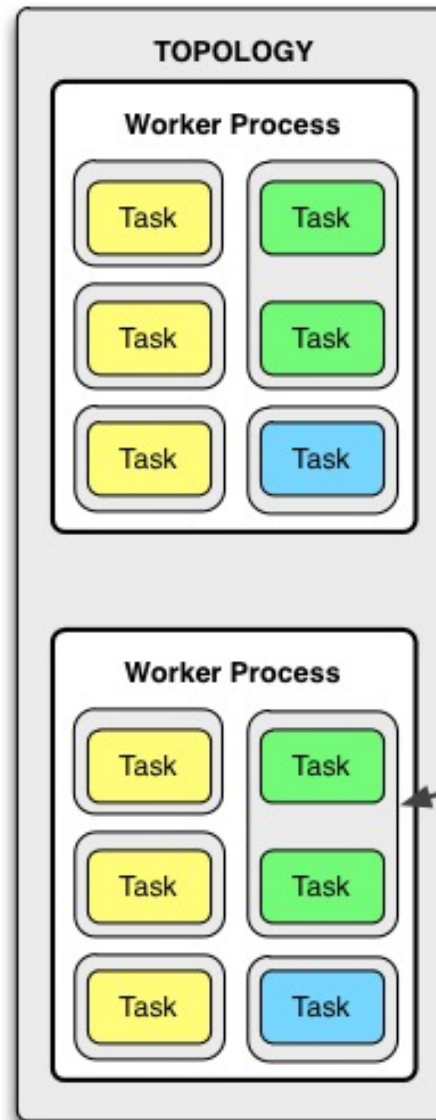


Parallelism in Storm

- A topology is ran by 3 entities
 - Worker (processes)
 - Executor (threads)
 - Tasks
- When creating a topology, it's possible to change the number of these entities
- Workers:
 - How many workers processes to create for the topology
 - `org.apache.storm.Config.setNumWorkers(...)`
- Executors
 - How many executors to spawn per Spout/Bolt
 - Parameter *parallelismHint*
- Tasks
 - How many tasks per Spout/Bolt
 - Method `setNumTasks(...)`
 - If not set, Storm will create 1 task per executor

2 executors

- Workers set to 2
- Blue Spout
 - ParallelismHint : 2
 - SetNumTasks : 1
- Green Bolt
 - ParallelismHint : 2
 - Tasks : 4
- Yellow Bolt
 - ParallelismHint : 6
 - SetNumTasks : 1



Case Study : Apache Storm

Writing topologies

Writing Spouts

- A spout should
 - Extends BaseRichSpout
 - Indicates what tuples it will produce
 - Provides a mechanism to emit new tuples
- There can be many instances of the same Spout
 - Can use private attributes to save state
- Spouts are started by the Storm runtime

Tuples
└ "entier" : value

Writing Spouts

- Declaring output tuples
 - Must override *declareOutputFields*
 - Each field in the tuple has a String id, but no type is specified


```
@Override  
public void declareOutputFields(OutputFieldsDeclarer declarer) {  
    declarer.declare(new Fields(...fields:"entier"));  
}
```

on the default stream

- Possible to declare multiple fields

```
@Override  
public void declareOutputFields(OutputFieldsDeclarer declarer) {  
    declarer.declare(new Fields(...fields:"word", "count"));  
}
```

Writing Spouts



- A spout can declare Streams
 - Useful for partitioning output data
 - By default id is “default”

```
@Override  
public void declareOutputFields(OutputFieldsDeclarer declarer) {  
    declarer.declareStream(streamId:"entier", new Fields(...fields:"entier"));  
}
```

Writing Spouts

- When a task is created by Storm runtime, the open method will be called
 - Gives a reference to the SpoutOutputCollector to emit
 - Useful for instancing private variables

```
@Override  
public void open(Map conf, TopologyContext context, SpoutOutputCollector collector) {  
    _collector = collector;  
    _rand = new Random();  
}
```

Writing Spouts

- Storm will request the Spout to emit new Tuples with method *nextTuple*
 - Should be non blocking
 - Will be called most of the time without pause

```
@Override
public void nextTuple() {
    Utils.sleep(millis:10);
    String[] sentences = new String[]{ "the cow jumped over the moon", "an apple a
    |   "four score and seven years ago", "snow white and the seven dwarfs", "i am
    String sentence = sentences[_rand.nextInt(sentences.length)];
    _collector.emit(new Values(sentence));
}
```

Send on Stream

Writing Bolts

- A Bolt should
 - Extends BaseBasicBolt
 - Indicates what tuples it will produce
 - Provides a mechanism to process incoming tuples
- There can be many instances of the same Bolt
 - Can use private attributes to save state
- Bolts are started by the Storm runtime

Writing Bolts

- Declaring output fields
 - Just like in Spouts, it can declare Fields and Streams

```
@Override  
public void declareOutputFields(OutputFieldsDeclarer declarer) {  
    declarer.declare(new Fields(...fields:"word", "count"));  
}
```

Writing Bolts

- When a task is created by Storm runtime, the *prepare* method will be called
 - Does NOT give a reference to collector
 - Useful for instancing private variables

```
@Override  
public void prepare(Map stormConf, TopologyContext context) {  
  
}
```

Writing Bolts

- When a tuple is available, the Storm runtime calls *execute*
 - Pass the Tuple as argument
 - Pass the collector also
- Fields in Tuple can be accessed by String id
 - Need to know the type to call the corresponding getXXXByField

```
public void execute(Tuple tuple, BasicOutputCollector basicOutputCollector) {  
    Integer v = tuple.getIntegerByField(field:"entier");  
    basicOutputCollector.emit(new Values(v*2));  
}
```

Writing topologies

- A topology links Spouts and Bolts with Streams
- Created using a *TopologyBuilder*
- Describe the topology
 - Each component has a String id used for linking
 - Order is destination, source
 - Must specify the type of grouping
 - Optionally specify the parallelism level and number of tasks

```
TopologyBuilder builder = new TopologyBuilder();  
builder.setSpout(id:"spout", new RandomSentenceSpout());  
builder.setBolt(id:"split", new SplitSentence()).shuffleGrouping(componentId:"spout");  
builder.setBolt(id:"count", new WordCount()).fieldsGrouping(componentId:"split", new Fields(...fields:"word"));
```



Wrapping-up everything

- Write startup code in a *main* method
 1. Create TopologyBuilder
 2. Describe your topology
 3. Create an instance of *org.apache.storm.Config*
 4. Write start-up/submission code
- Build with Maven

Wrapping-up everything

- Execution like with Hadoop
 - Submit jar with the *storm* command
- Execution can be local or cluster-based
 - Specified in the source code

✓ storm jar ---

```
if (args != null && args.length > 0) {  
    StormSubmitter.submitTopologyWithProgressBar(args[0], conf, builder.createTopology());  
} else {  
    LocalCluster cluster = new LocalCluster();  
    cluster.submitTopology(topologyName:"word-count", conf, builder.createTopology());  
    Thread.sleep(10000);  
    cluster.shutdown();  
}
```

Windowing Support

- Storm supports both sliding and tumbling windows
 - Implemented in
 - interface *IWindowedBolt*
 - Class *BaseWindowedBolt*
 - When instancing Bolt, specify
 - Length of window
 - Sliding interval in number of tuples
 - Length of window can be time-based or number-based
 - If time-based, late tuples are logged by default
-