

Heavy virtualization

Outline

- Virtualization techniques and hypervisor
- Virtualization without architectural support
 - Xen & VMWare (good old days...)
- The Popek/Goldberg Theorem
- Virtualization with architectural support
 -


Introduction - 1

- Virtualizations boils down to separating a service from the underlying layer (hardware)
- Old technique pioneered in the 60s
 - Used in mainframes
 - But also directly in operating systems



Introduction - 2

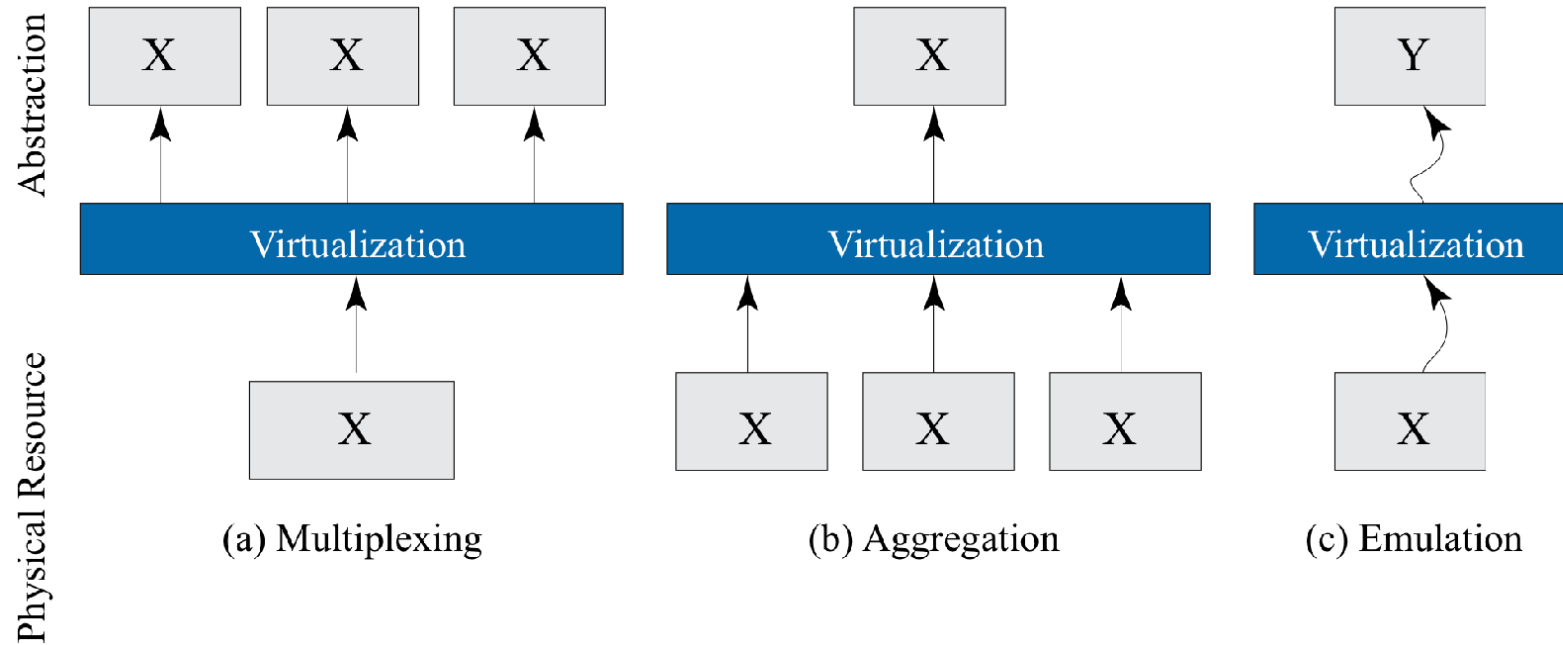
- Two fundamental principles
 - Layering : presentation of a single abstraction
 - Enforced modularity :
 - The exposed higher-level interface is the same as the lower-level one
 - The higher-level software cannot escape the modularity
- Virtualization is not limited to Virtual Machines
- Example 1 : Redundant Array of Inexpensive Disk (RAID)
 - Aggregation of multiple disk as a single (virtual) disk



Introduction- 3

- Example 2 : Memory management in Operating Systems
- Each process has its own (complete) address space
 - Physical memory is shared among them
 - No direct access to RAM, always go through a Memory Management Unit (MMU)
 - More on that later...

Virtualization techniques



- Multiplexing : OS with physical memory
- Aggregation : RAID
- Emulation : Emulate a disk in RAM (e.g. /proc in Linux)

Hypervisor

Key features (Popek and Goldberg, 1974)

- A software providing virtualization service is called a *Hypervisor*
 - Virtualized machines are sometimes called guests
- Must have 3 key features (Popek and al.)
 - Equivalent execution
 - Software running in a virtualized environment should have an identical execution to running in a native one, except for resource usage and timing
 - Safety ↳ slower
 - Virtual machines are isolated from each other as well as from the hypervisor
 - Performance
 - VM should suffer only from minor performance decrease. Most instructions should be executed directly on the physical CPU.
 - Difference between hypervisors and simulators

Instructions and CPU

- A program is made of instructions
 - Executed by the CPU
- Not all instructions are equals
 - Some instructions are free to use (computation)
 - Some instructions are only accessible to the Kernel (management of resources)
- 2 set of instructions
 - Unprivileged : can be called by any process
 - Privileged : only callable by the kernel
- Enforced by the CPU itself

x86 rings

- Standard x86 architecture has 4 privilege levels
 - Rings 0-3
- Each ring restricts callable instructions
 - When calling a forbidden instruction, CPU trap
- Processes are executed in Ring 3
- OS is executed in Ring 0
 - Has access to all instructions
- Others rings are unused.

System calls - 1

- How can an unprivileged process execute privileged code
 - Need a mechanism to temporary change ring level
 - It's the system calls
- The OS provides a list of system calls
 - Specified with numbers
- Different ways to call them
- Software interrupt (*legacy*)
 - Unprivileged process write the system call number in EAX CPU register
 - Process then call software interrupt 0x80
 - The OS is executed and looks into EAX to continue

System calls - 2

- On x86 in 32/64 bits modes, 0x80 is deprecated.
- 32 bits mode: use of 2 new instructions + EAX
 - Syntenter : quickly switch to ring 0
 - Systexit : quickly switch back to ring 3
- In 64 bits mode
 - Syscall
 - Sysret

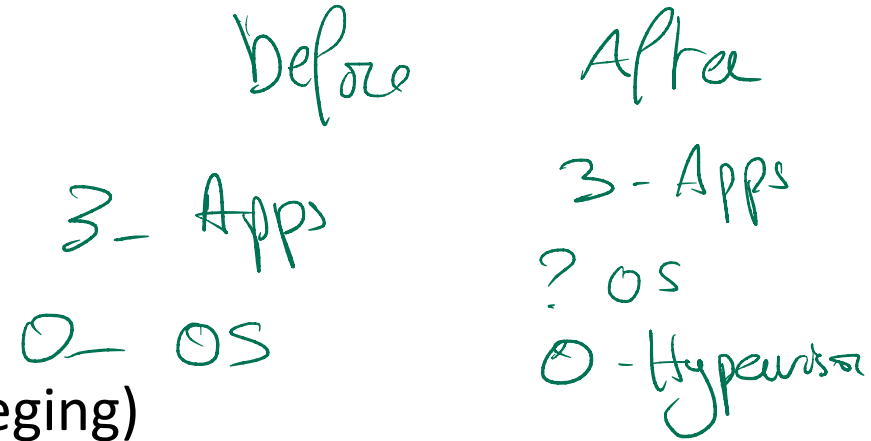
↪ register

Trap


- A form of hardware interrupt
- Raised by the CPU when something is wrong
 - Invalid memory access
- At boot time, the OS registers methods (memory address) to be executed in case of trap
- Allows for handling of errors
 - Ex : invalid memory access by a process => trap to OS => kill the process (segfault anyone?)

Trap-and-emulate virtualization - 1

- Allows for virtualization on “classical” x86 CPUs
- Slightly changes ring levels
 - Hypervisor in Ring 0 (all access)
 - OS moved to less privileged rings (aka Ring deprivileging)
- What happens when guest OS calls a ring 0 instruction ?
 - CPU error and trap
 - Hypervisor is executed and handles the instruction

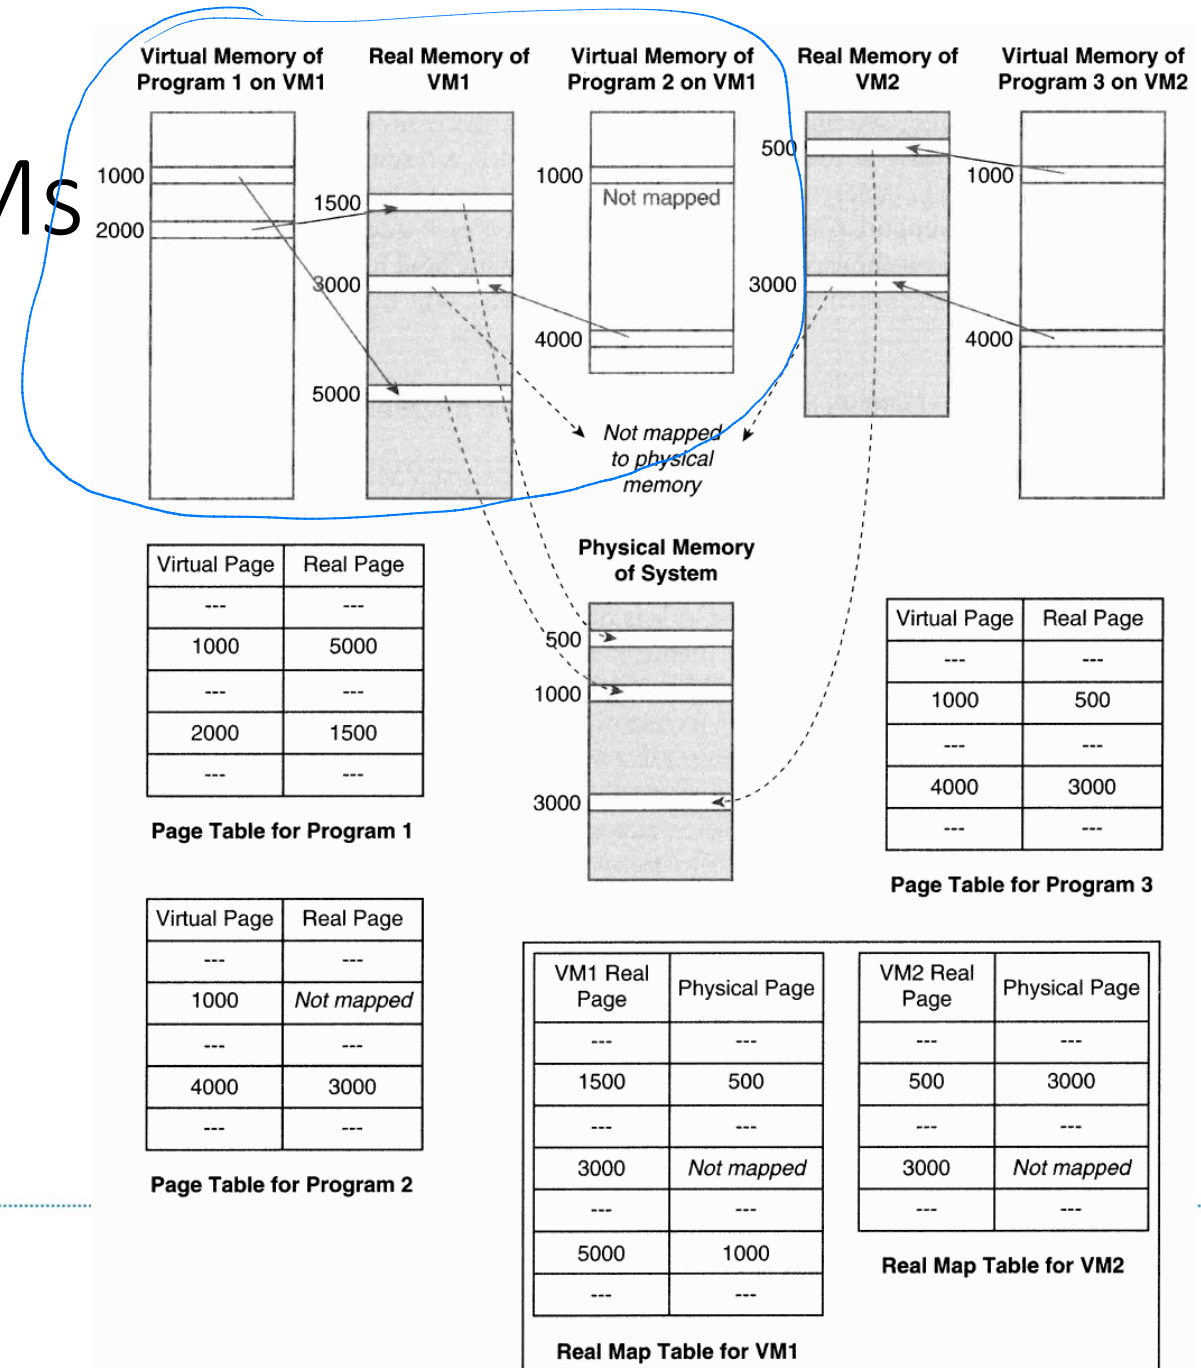


Classical Virtual Memory



Virtual Memory with VMs

- 2 indirections level
 - Guest Program -> Guest OS
 - Guest OS -> host

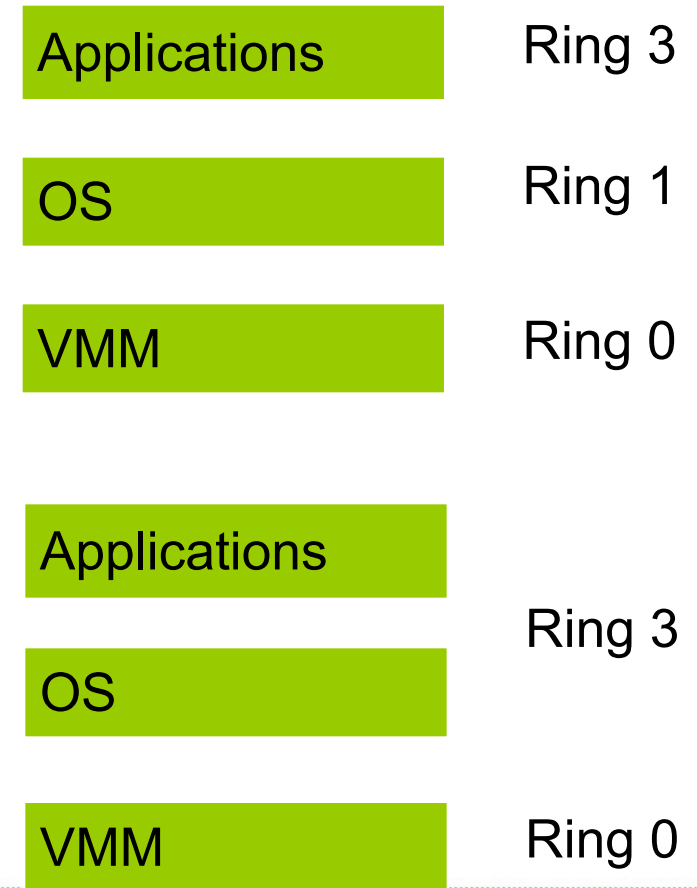


Virtualization of CPU, RAM and I/O

- CPU and RAM virtualization
 - RAM : spatial multiplexing
 - CPU : temporal multiplexing
- I/O virtualization done via emulation
 - Key advantage : portability (guest sees the same virtual hardware on different hosts)
 - Better hardware support now (see later)
 - Why not direct access to device ?

Trap-and-emulate virtualization – 2

- 2 different models



Discussion

- Rings 0/3 :
 - Called Ring Compression
 - OS and user programs are in the same ring level
 - So no protection between them (bad)
- Rings 0/1/3 :
 - OS is protected from user programs
 - Hypervisor is protected from OS
- So it's a good solution... or not
 - On x86 some instructions fail silently when called from the wrong ring

Getting around x86 limitations without hardware support

- How to force trap on some instructions ?
- Solution 1: paravirtualisation
 - Modify the guest OS
 - Replace non-trappable system calls with calls to hypervisor
 - Not too hard if source code is available
 - Example : Xen on Linux
 - Still exists today (guest additions...)
- Solution 2: binary translation
 - Execute the guest OS on an interpreter
 - Modify it dynamically
 - Dynamic and adaptive translation
 - Example : VMWare

Paravirtualization

Xen

Xen

- Xen is an hypervisor that uses paravirtualization
- Guest OS has to be modified...
 - XenoLinux kernel
- ... but not the applications
- Memory management
 - Guest OS are responsible for memory management
 - Xen keeps track of page directory and page tables
- When a guest OS creates a new Page table
 - Register it with Xen
 - All following updates will be validated by Xen
 - Tables are read only for guests

Xen - 2

- Guest OS runs in ring 1
- Paravirtualization of privileged instructions
- Xen handles interrupt with specific handlers
 - E.g page fault
- I/O
 - Pass interrupts to the corresponding guest
 - Memory transfer between device and guest using shared memory
 - Handled by Xen

Xen – Performance

- Old benchmarks (2003)
- Performance will vary greatly with workload
 - CPU vs Memory vs IO
- Spec CPU
 - Same performance as native
- Linux kernel 2.4 compilation (7% CPU, 93% I/O)
 - 3% overhead
- OS bench
 - Fork : 198 μ s (Xen) vs 110 μ s (Linux native)
 - Context switch : up to 3 μ s overhead

Binary translation

VMware

VMWare

- Founded in 1998
- Interpretation/compilation of guest OS code
- Idea : the hypervisor keeps track of important register/structures
 - Use of **shadow** structures
- But hard for memory area with sensitive data structures
 - E.g page tables
 - Could be modified by any memory access from guest
 - Need some protection

Binary translation

- Instructions from guest are grouped into Translation Units
- Computational instructions are not translated
- Jump/Branch instruction have to be modified
 - Jump addresses are not valid anymore
- No optimization performed

VMWare Performance

- VMPlayer in software mode
- SPECint 2000
 - Integer computation
 - Between 0 and 9 % overhead
- SPECjbb 2005
 - Server side java
 - 0-2% overhead
- Apache ab benchmark
- HTTP server performance
 - Between 25 and 65% overhead

Table 1.1: Virtual Hardware of early VMware Workstation [45]

	Virtual Hardware (front-end)	Back-end
Multiplexed	1 virtual x86-32 CPU	Scheduled by the host operating system with one or more x86 CPUs
	Up to 512 MB of contiguous DRAM	Allocated and managed by the host OS (page-by-page)
Emulated	PCI Bus	Fully emulated compliant PCI bus with B/D/F addressing for all virtual motherboard and slot devices
	4 x 4IDE disks 7 x Buslogic SCSI Disks	Either virtual disks (stored as files) or direct access to a given raw device
	1 x IDE CD-ROM	ISO image or real CD-ROM
	2 x 1.44 MB floppy drives	Physical floppy or floppy image
	1 x VGA/SVGA graphics card	Appears as a Window or in full-screen mode
	2 x serial ports COM1 and COM2	Connect to Host serial port or a file
	1 x printer (LPT)	Can connect to host LPT port
	1 x keyboard (104-key) and mouse	Fully emulated
	AMD PCnet NIC (AM79C970A)	Via virtual switch of the host

The Popek/Goldberg Theorem

Is an ISA Virtualizable?

- What are the requirements to virtualize an ISA?
 - Necessary and sufficient
- “Formal Requirements for Virtualizable Third-Generation Architectures”, Popek and Goldberg, in Communications of the ACM, 1974
- The theorem states that if hypothesis met, an OS running on hardware can also run inside a virtual machine

Model

- One processor with two modes of execution
 - User-level and supervisor
- Virtual memory implemented with segments (not pages)
 - Virtual address $x \in [0, L]$ mapped to segment with offset B (i.e. $[B, B+L]$)
- Physical memory is contiguous
- Processor state (called processor status word – psw) consists of
 - Execution level (user, supervisor)
 - Segment register (B, L)
 - Current Program counter

Model - 2

- Trap
 - Saves PSW to a well location
 - Load into PSW values from another well-known location
 - What is the purpose... ? *Manage how then come back to normal flow*
- ISA support instructions to
 - Store/Load PSW from arbitrary virtual location
 - What is the purpose... ? *→ switch between VTT*

Classification of instruction

- Simple privileged/non-privileged distinction not enough
- Control-sensitive instruction
 - Can update the system state
- Behavior-sensitive instruction
 - its semantics depend on the actual values set in the system state
- Innocuous instruction
 - All the others

Elements of proof

- Hypervisor is only software in supervisor mode
- Each VM has a fixed and contiguous part of physical memory
 - Segment hypothesis
 - Enforce isolation
- Hypervisor has reserved physical memory, separated from VMs
- Hypervisor maintain copies of the psw of each VM

Hardware assisted virtualization

Intel and AMD

- x86 ISA common to Intel and AMD processors
- Intel Virtualization Technology (Intel VT)
- AMD Virtualization (AMD-V)
- First introduced in 2006
 - New features added regularly
- Objective
 - Eliminate the need for paravirtualization and binary translation

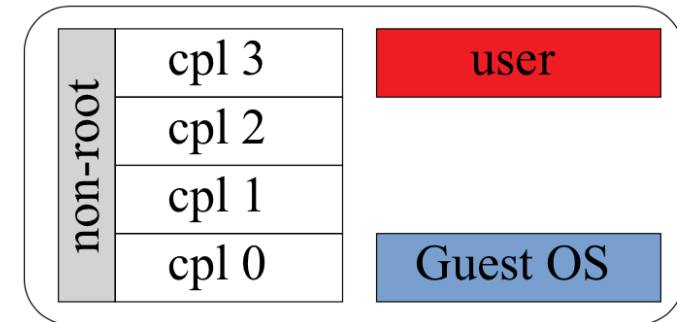
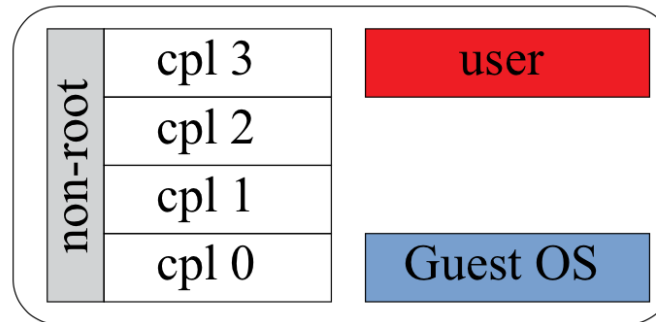
Intel VT-x

- Do not change the semantics of instructions
 - Don't remove 17 sensitive instructions
- Duplicates the visible state of the processor
 - Add a new mode of execution : root mode
- Hypervisor and OS run in root mode
- VMs run in non-root mode
- Properties
 - Transition between root and non-root mode is atomic
 - Each mode still relies on security rings


Host

VM 1

VM 2

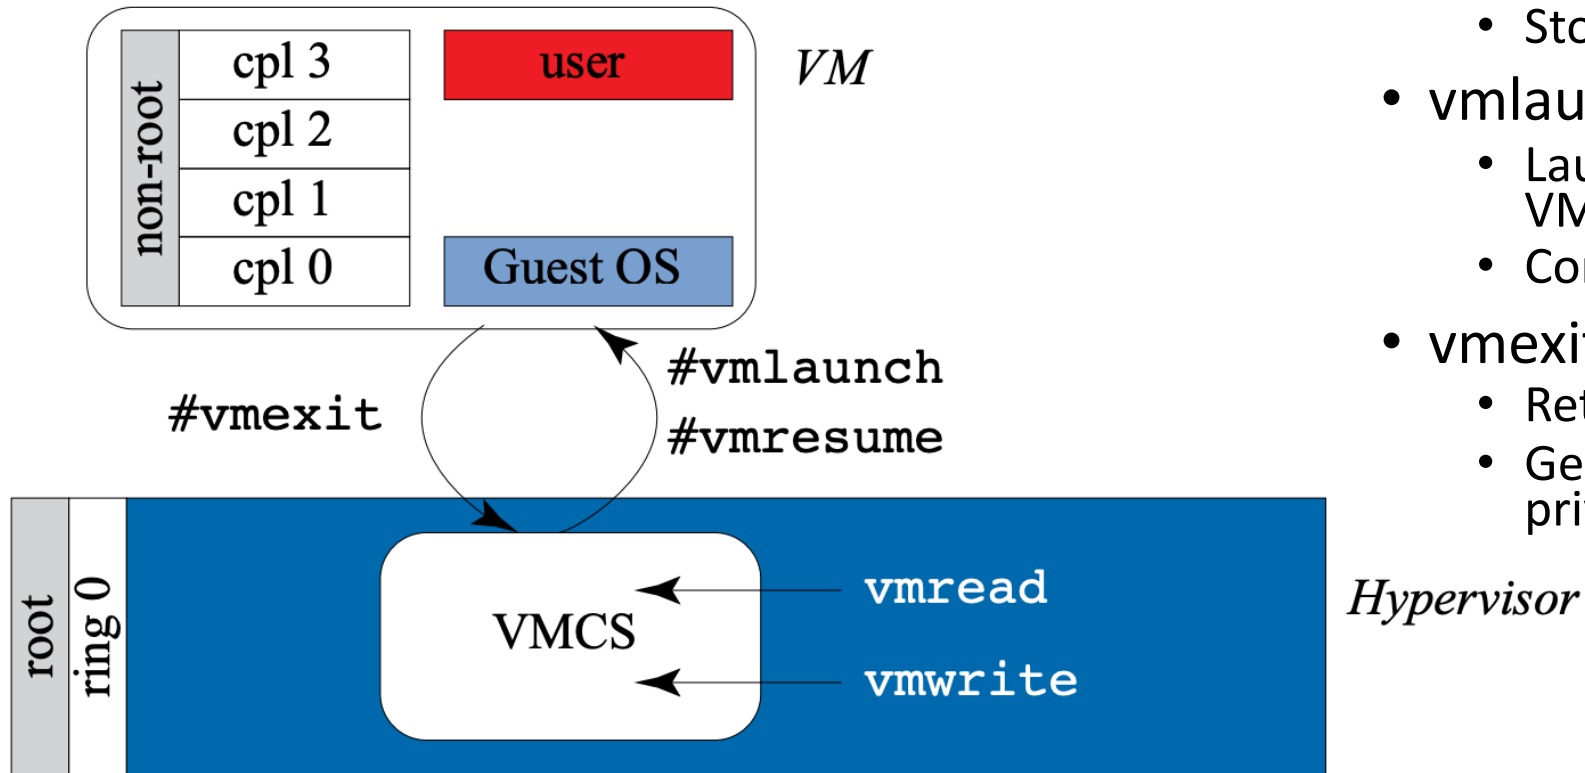


Standard
no VM



Hypervisor and/or Host Operating System

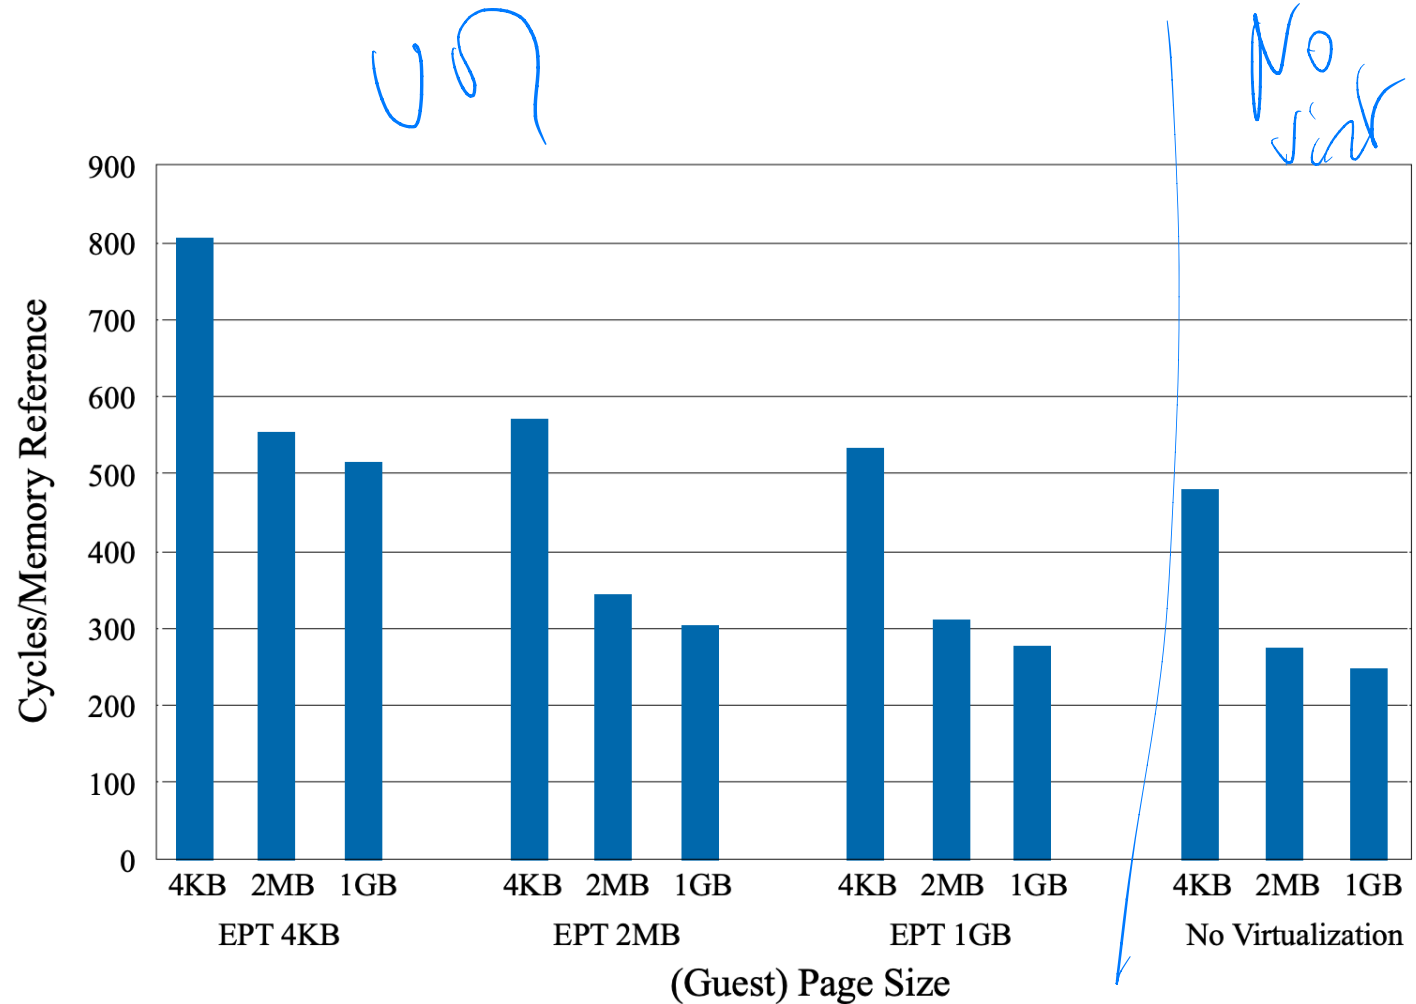
VT-x transitions




- Virtual Machine Control Structure (VMCS)
 - Stores the state of the VM
- vmlaunch/vmresume
 - Launches/resume a VM managed by the VMCS
 - Control is transferred to VM
- vmexit
 - Returns control to hypervisor
 - Generated when calling root-mode-privileged instructions, interrupts...

MMU Virtualization


- Extended Page Tables (EPT)
 - Available on Intel and AMD
- Add a second page table structure
 - Map guest-physical to host-physical address



Virtual I/O with hardware support



(a) Emulation/paravirtualization



(b) Direct device assignment

- Allow guest direct access to device (passthrough)
 - Not scalable
 - Security issue with DMA

Virtual I/O with hardware support

- 2 mechanisms
 - IOMMU (I/O Memory management unit)
 - SRIOV (Single Root I/O Virtualization)
- IOMMU provides security
 - Remapping of DMA
 - Instead of DMA to physical address, DMA to I/O virtual address
- SRIOV extends PCIe
 - Support for devices that self-virtualize
 - An SRIOV-capable I/O device can present multiples instances of itself
 - E.g : a GPU, a NIC